

Hallucination & Bias Detection Report

Deliverable 2 — Bias Assessment Section (Final, Reconciled Version)

Source reviewed: Oluchi Loveth Agbo — Version C (full AI output and verification matrix) | Prepared for: Alishba

This version reconciles two independent reviews of the same Version C output. Where the two reviews disagreed, the underlying text was re-checked directly (including a word count of each theme paragraph and a re-check of the verification matrix) to determine which reading the evidence supports.

1. Four-Question Bias Assessment

Are multiple perspectives considered?

Yes, at the surface level — confirmed by direct word count. The overreliance/disagreement theme (≈ 70 words) is in fact the longest of the three, not the shortest, so the concern is not under-covered. The residual bias is not length but sequencing: the concern is stated first and then closed out by a rebuttal (“However, adoption research ... frames this as an implementation and training gap”), so the AI-favorable interpretation gets the last word in that paragraph.

Is the language neutral?

Mostly. One line overstates its source: OECD's 2026 report is described as “confirming” adaptive platforms are scaling, but per the verification matrix the report's actual framing is more cautious and centers a gap between digital expansion and real transformation. This is a wording precision issue more than an emotional-language issue.

Are assumptions clearly identified?

No, on both reviews' agreement. The output does not flag that the headline effect size (0.73–1.3 SD) comes from one crossover trial of 194 undergraduate physics students at a single university, nor does it note that findings from Western/OECD contexts are being generalized to “classroom learning” broadly.

Is evidence balanced?

Balanced in volume, not fully balanced in resolution. All three themes get comparable space and each cites verified, credible sources. However, the paragraph that raises the strongest counter-evidence (overreliance) is also the one paragraph that talks itself back down by its final sentence.

2. Biases Identified

Confirmation Bias — Low, with one caveat

What it is: A tendency to favor evidence supporting AI benefits and downplay or soften contradicting evidence.

Where it appears: Not in the overall structure (the concern paragraph is the longest section, so the response is not simply ignoring criticism). It appears narrowly in the final sentence of the overreliance theme, where the faculty/researcher concern is reframed as a “training gap rather than an inherent flaw” — letting the AI-favorable interpretation close the paragraph.

How to correct: Present the RAND/CDT reframing as a co-equal, separately labeled viewpoint (e.g., “Counter-interpretation:”) rather than the concluding sentence of the concern paragraph, so neither side is structurally positioned as the final word.

Selection Bias — Moderate

What it is: Reliance on one type of evidence (quantitative) while excluding others (qualitative).

Where it appears: All six sourced claims in the verification matrix (C-01–C-06) are RCT, survey, or systematic-review statistics. No classroom observation, teacher/student interview, or case-study evidence is included.

How to correct: Add at least one qualitative or mixed-methods source to capture classroom experience that percentages alone do not convey.

Cultural Bias — Moderate

What it is: A geographically narrow evidence base presented without acknowledging the limitation.

Where it appears: Every cited organization — Harvard (Kestin et al.), CDT, AAC&U/Elon, RAND, and OECD — is US-based or Western-majority. No source covers the Global South or non-Western education systems, and the output never notes this as a limitation.

How to correct: Add a limitation statement and, where possible, cite research from a wider range of regions before generalizing to “classroom learning” as a whole.

3. Additional Finding: Unverified Source Feeding the Bias

The overreliance theme cites “adoption research (RAND, CDT)” to support the training-gap reinterpretation. CDT is verified in the matrix (C-04); **RAND is not** — no RAND source appears anywhere in the verification matrix (C-01–C-06). This means one of the two sources used to soften the strongest critical finding in the entire output is uncorroborated. This is both a hallucination-check gap and a bias-relevant finding: an unverified citation is doing rhetorical work specifically in the direction of downplaying the concern.

Recommendation: either locate and verify a specific RAND publication supporting this claim, or remove the RAND attribution and rely only on the verified CDT framing.

4. Summary Recommendation

Version C is well-sourced and, by word count, more balanced than an initial read suggested — the concern section is not shortchanged in length. The remaining issues are narrower than “overall confirmation bias”: (1) reposition the RAND/CDT reframing as a parallel viewpoint rather than the paragraph’s closing rebuttal, (2) verify or drop the RAND citation, (3) add qualitative evidence, and (4) add a brief caveat on the Western/OECD-centric source base before generalizing globally.